

In [1]:

```
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

Using TensorFlow backend.

In [0]:

```
import matplotlib.pyplot as plt
import numpy as np
import time
import pandas as pd

from keras.models import Sequential
from keras.layers import Dense, Activation

from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout
```

In [0]:

```
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [6]:

```
(X_train, y_train), (X_test, y_test) = mnist.load_data()
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>
11493376/11490434 [=====] - 0s 0us/step

In [0]:

```
X_train = X_train/255
X_test = X_test/255
```

In [0]:

```
Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)
```

In [0]:

```
output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

2 Layers

In [28]:

```
model_drop = Sequential()
```

```

model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(
mean=0.0, stddev=0.062, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(80, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()

```

Layer (type)	Output Shape	Param #
dense_12 (Dense)	(None, 364)	285740
batch_normalization_9 (Batch Normalization)	(None, 364)	1456
dropout_9 (Dropout)	(None, 364)	0
dense_13 (Dense)	(None, 80)	29200
batch_normalization_10 (Batch Normalization)	(None, 80)	320
dropout_10 (Dropout)	(None, 80)	0
dense_14 (Dense)	(None, 10)	810
Total params: 317,526		
Trainable params: 316,638		
Non-trainable params: 888		

In [29]:

```

model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 9s 150us/step - loss: 0.5568 - acc: 0.8326 - val_loss: 0.1769 - val_acc: 0.9444
Epoch 2/20
60000/60000 [=====] - 8s 126us/step - loss: 0.2668 - acc: 0.9217 - val_loss: 0.1293 - val_acc: 0.9592
Epoch 3/20
60000/60000 [=====] - 7s 120us/step - loss: 0.2066 - acc: 0.9394 - val_loss: 0.1076 - val_acc: 0.9659
Epoch 4/20
60000/60000 [=====] - 7s 118us/step - loss: 0.1762 - acc: 0.9473 - val_loss: 0.0936 - val_acc: 0.9715
Epoch 5/20
60000/60000 [=====] - 7s 122us/step - loss: 0.1520 - acc: 0.9543 - val_loss: 0.0861 - val_acc: 0.9743
Epoch 6/20
60000/60000 [=====] - 7s 121us/step - loss: 0.1373 - acc: 0.9595 - val_loss: 0.0832 - val_acc: 0.9748
Epoch 7/20
60000/60000 [=====] - 7s 119us/step - loss: 0.1272 - acc: 0.9623 - val_loss: 0.0784 - val_acc: 0.9750
Epoch 8/20
60000/60000 [=====] - 7s 121us/step - loss: 0.1188 - acc: 0.9640 - val_loss: 0.0757 - val_acc: 0.9768
Epoch 9/20
60000/60000 [=====] - 7s 121us/step - loss: 0.1110 - acc: 0.9666 - val_loss: 0.0730 - val_acc: 0.9776

```

```

00000/00000 [=====] - 7s 121us/step - loss: 0.1110 - acc: 0.9000 - val_loss: 0
.0740 - val_acc: 0.9784
Epoch 10/20
60000/60000 [=====] - 7s 123us/step - loss: 0.1060 - acc: 0.9685 - val_loss: 0
.0678 - val_acc: 0.9789
Epoch 11/20
60000/60000 [=====] - 7s 121us/step - loss: 0.0964 - acc: 0.9710 - val_loss: 0
.0703 - val_acc: 0.9786
Epoch 12/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0904 - acc: 0.9725 - val_loss: 0
.0686 - val_acc: 0.9797
Epoch 13/20
60000/60000 [=====] - 7s 121us/step - loss: 0.0879 - acc: 0.9732 - val_loss: 0
.0680 - val_acc: 0.9791
Epoch 14/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0859 - acc: 0.9739 - val_loss: 0
.0672 - val_acc: 0.9790
Epoch 15/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0787 - acc: 0.9760 - val_loss: 0
.0607 - val_acc: 0.9815
Epoch 16/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0773 - acc: 0.9765 - val_loss: 0
.0675 - val_acc: 0.9803
Epoch 17/20
60000/60000 [=====] - 7s 121us/step - loss: 0.0765 - acc: 0.9759 - val_loss: 0
.0581 - val_acc: 0.9822
Epoch 18/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0729 - acc: 0.9774 - val_loss: 0
.0615 - val_acc: 0.9819
Epoch 19/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0704 - acc: 0.9784 - val_loss: 0
.0628 - val_acc: 0.9821
Epoch 20/20
60000/60000 [=====] - 7s 125us/step - loss: 0.0669 - acc: 0.9795 - val_loss: 0
.0631 - val_acc: 0.9810
Test score: 0.06307864126095082
Test accuracy: 0.981

```

In [30]:


```

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
plt.show()

```



In [31]:

```

w_after = model_drop.get_weights()

```

```
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
```

```
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')


plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is deprecated and is a private function. Do not use.

```
kde_data = remove_na(group_data)
```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is deprecated and is a private function. Do not use.

```
violin_data = remove_na(group_data)
```



3 Layers

In [32]:

```
model_drop = Sequential()

model_drop.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(
mean=0.0, stddev=0.062, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.088, se
ed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, se
ed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))
```

```
model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_15 (Dense)	(None, 512)	401920
batch_normalization_11 (Batch Normalization)	(None, 512)	2048
dropout_11 (Dropout)	(None, 512)	0
dense_16 (Dense)	(None, 256)	131328
batch_normalization_12 (Batch Normalization)	(None, 256)	1024
dropout_12 (Dropout)	(None, 256)	0
dense_17 (Dense)	(None, 128)	32896
batch_normalization_13 (Batch Normalization)	(None, 128)	512
dropout_13 (Dropout)	(None, 128)	0
dense_18 (Dense)	(None, 10)	1290
Total params: 571,018		
Trainable params: 569,226		
Non-trainable params: 1,792		

In [34]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 11s 187us/step - loss: 0.6390 - acc: 0.8050 - val_loss: 0.1855 - val_acc: 0.9417
Epoch 2/20
60000/60000 [=====] - 10s 161us/step - loss: 0.2754 - acc: 0.9183 - val_loss: 0.1281 - val_acc: 0.9615
Epoch 3/20
60000/60000 [=====] - 9s 157us/step - loss: 0.2049 - acc: 0.9396 - val_loss: 0.1085 - val_acc: 0.9662
Epoch 4/20
60000/60000 [=====] - 9s 152us/step - loss: 0.1734 - acc: 0.9488 - val_loss: 0.0939 - val_acc: 0.9721
Epoch 5/20
60000/60000 [=====] - 9s 155us/step - loss: 0.1482 - acc: 0.9563 - val_loss: 0.0913 - val_acc: 0.9719
Epoch 6/20
60000/60000 [=====] - 9s 157us/step - loss: 0.1360 - acc: 0.9596 - val_loss: 0.0815 - val_acc: 0.9766
Epoch 7/20
60000/60000 [=====] - 9s 156us/step - loss: 0.1189 - acc: 0.9641 - val_loss: 0.0791 - val_acc: 0.9766
Epoch 8/20
60000/60000 [=====] - 9s 154us/step - loss: 0.1117 - acc: 0.9663 - val_loss: 0.0752 - val_acc: 0.9761
Epoch 9/20
60000/60000 [=====] - 9s 153us/step - loss: 0.1044 - acc: 0.9686 - val_loss: 0.0675 - val_acc: 0.9802
Epoch 10/20
60000/60000 [=====] - 9s 153us/step - loss: 0.0969 - acc: 0.9709 - val_loss: 0.0687 - val_acc: 0.9797
Epoch 11/20
60000/60000 [=====] - 10s 159us/step - loss: 0.0921 - acc: 0.9724 - val_loss: 0.0672 - val_acc: 0.9807
```

```

0.0072 - val_acc: 0.9807
Epoch 12/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0883 - acc: 0.9741 - val_loss: 0
.0707 - val_acc: 0.9804
Epoch 13/20
60000/60000 [=====] - 10s 160us/step - loss: 0.0829 - acc: 0.9743 - val_loss:
0.0660 - val_acc: 0.9818
Epoch 14/20
60000/60000 [=====] - 10s 163us/step - loss: 0.0810 - acc: 0.9755 - val_loss:
0.0650 - val_acc: 0.9802
Epoch 15/20
60000/60000 [=====] - 10s 164us/step - loss: 0.0754 - acc: 0.9772 - val_loss:
0.0657 - val_acc: 0.9831
Epoch 16/20
60000/60000 [=====] - 10s 161us/step - loss: 0.0748 - acc: 0.9772 - val_loss:
0.0608 - val_acc: 0.9812
Epoch 17/20
60000/60000 [=====] - 10s 161us/step - loss: 0.0725 - acc: 0.9773 - val_loss:
0.0662 - val_acc: 0.9813
Epoch 18/20
60000/60000 [=====] - 10s 164us/step - loss: 0.0678 - acc: 0.9797 - val_loss:
0.0615 - val_acc: 0.9832
Epoch 19/20
60000/60000 [=====] - 10s 162us/step - loss: 0.0635 - acc: 0.9805 - val_loss:
0.0588 - val_acc: 0.9837
Epoch 20/20
60000/60000 [=====] - 10s 161us/step - loss: 0.0663 - acc: 0.9803 - val_loss:
0.0539 - val_acc: 0.9849
Test score: 0.053948233158088985
Test accuracy: 0.9849

```

In [35]:


```

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
plt.show()

```



In [36]:

```

w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()

```

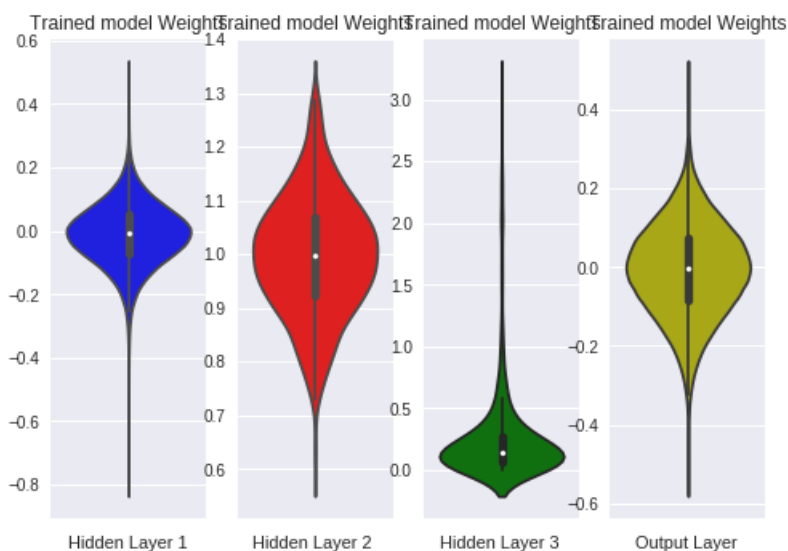
```
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is deprecated and is a private function. Do not use.
 kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is deprecated and is a private function. Do not use.
 violin_data = remove_na(group_data)



5 Layers

In [37]:

```
model_drop = Sequential()

model_drop.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(
mean=0.0, stddev=0.062, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.088, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.176, seed=None)))
model_drop.add(BatchNormalization())
```

```

model_drop.add(Dropout(0.5))

model_drop.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.25, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()

```

Layer (type)	Output Shape	Param #
dense_19 (Dense)	(None, 512)	401920
batch_normalization_14 (Batch Normalization)	(None, 512)	2048
dropout_14 (Dropout)	(None, 512)	0
dense_20 (Dense)	(None, 256)	131328
batch_normalization_15 (Batch Normalization)	(None, 256)	1024
dropout_15 (Dropout)	(None, 256)	0
dense_21 (Dense)	(None, 128)	32896
batch_normalization_16 (Batch Normalization)	(None, 128)	512
dropout_16 (Dropout)	(None, 128)	0
dense_22 (Dense)	(None, 64)	8256
batch_normalization_17 (Batch Normalization)	(None, 64)	256
dropout_17 (Dropout)	(None, 64)	0
dense_23 (Dense)	(None, 32)	2080
batch_normalization_18 (Batch Normalization)	(None, 32)	128
dropout_18 (Dropout)	(None, 32)	0
dense_24 (Dense)	(None, 10)	330
Total params: 580,778		
Trainable params: 578,794		
Non-trainable params: 1,984		

In [38]:

```

model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 17s 276us/step - loss: 1.7274 - acc: 0.4389 - val_loss: 0.4358 - val_acc: 0.8914
Epoch 2/20
60000/60000 [=====] - 13s 223us/step - loss: 0.7210 - acc: 0.7873 - val_loss: 0.2211 - val_acc: 0.9370
Epoch 3/20
60000/60000 [=====] - 13s 224us/step - loss: 0.4643 - acc: 0.8757 - val_loss: 0.1661 - val_acc: 0.9541
Epoch 4/20
60000/60000 [=====] - 13s 225us/step - loss: 0.3627 - acc: 0.9081 - val_loss: 0.1441 - val_acc: 0.9617

```



```

0.1441 - val_acc: 0.9611
Epoch 5/20
60000/60000 [=====] - 14s 226us/step - loss: 0.3083 - acc: 0.9257 - val_loss:
0.1298 - val_acc: 0.9660
Epoch 6/20
60000/60000 [=====] - 13s 224us/step - loss: 0.2693 - acc: 0.9365 - val_loss:
0.1182 - val_acc: 0.9704
Epoch 7/20
60000/60000 [=====] - 14s 228us/step - loss: 0.2408 - acc: 0.9446 - val_loss:
0.1114 - val_acc: 0.9727
Epoch 8/20
60000/60000 [=====] - 14s 228us/step - loss: 0.2231 - acc: 0.9479 - val_loss:
0.1131 - val_acc: 0.9718
Epoch 9/20
60000/60000 [=====] - 13s 224us/step - loss: 0.2105 - acc: 0.9513 - val_loss:
0.0968 - val_acc: 0.9766
Epoch 10/20
60000/60000 [=====] - 13s 225us/step - loss: 0.1956 - acc: 0.9556 - val_loss:
0.1040 - val_acc: 0.9752
Epoch 11/20
60000/60000 [=====] - 14s 228us/step - loss: 0.1809 - acc: 0.9590 - val_loss:
0.1028 - val_acc: 0.9751
Epoch 12/20
60000/60000 [=====] - 14s 226us/step - loss: 0.1753 - acc: 0.9599 - val_loss:
0.0966 - val_acc: 0.9765
Epoch 13/20
60000/60000 [=====] - 14s 226us/step - loss: 0.1688 - acc: 0.9613 - val_loss:
0.1037 - val_acc: 0.9763
Epoch 14/20
60000/60000 [=====] - 14s 225us/step - loss: 0.1655 - acc: 0.9627 - val_loss:
0.0891 - val_acc: 0.9779
Epoch 15/20
60000/60000 [=====] - 13s 225us/step - loss: 0.1561 - acc: 0.9643 - val_loss:
0.0812 - val_acc: 0.9804
Epoch 16/20
60000/60000 [=====] - 13s 225us/step - loss: 0.1459 - acc: 0.9664 - val_loss:
0.0858 - val_acc: 0.9804
Epoch 17/20
60000/60000 [=====] - 13s 224us/step - loss: 0.1492 - acc: 0.9665 - val_loss:
0.0876 - val_acc: 0.9797
Epoch 18/20
60000/60000 [=====] - 13s 223us/step - loss: 0.1385 - acc: 0.9684 - val_loss:
0.0797 - val_acc: 0.9799
Epoch 19/20
60000/60000 [=====] - 14s 226us/step - loss: 0.1365 - acc: 0.9700 - val_loss:
0.0815 - val_acc: 0.9804
Epoch 20/20
60000/60000 [=====] - 14s 227us/step - loss: 0.1300 - acc: 0.9714 - val_loss:
0.0824 - val_acc: 0.9810
Test score: 0.08238632999660912
Test accuracy: 0.981

```

In [39]:


```


fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
plt.show()

```





In [41]:

```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(2, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(2, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(2, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(2, 3, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(2, 3, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='b')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(2, 3, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is deprecated and is a private function. Do not use.

```
kde_data = remove_na(group_data)
```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is deprecated and is a private function. Do not use.

```
violin_data = remove_na(group_data)
```

